

Lecture 2: Spark DataFrame, Dataset, and ML Pipelines

Haiping Lu

<http://www.dcs.shef.ac.uk/~haiping>

COM6012 Scalable Machine Learning
Spring 2018

Week 2 Contents

- Spark Recap – Example on Cache
- Spark DataFrame & Dataset
- Spark MLlib & ML Pipelines
- GitHub Classroom & Quiz 1

Week 2 Contents

- **Spark Recap – Example on Cache**
- Spark DataFrame & Dataset
- Spark MLlib & ML Pipelines
- GitHub Classroom & Quiz 1

Apache Spark

- Fast and general cluster computing system, interoperable with Hadoop, included in all major distros
- Improves efficiency through:
 - In-memory computing primitives → Up to 100× faster (2-10× on disk)
 - General computation graphs
- Improves usability through:
 - Rich APIs in Scala, Java, Python → 2-5× less code
 - Interactive shell

Spark Model

- *Write programs in terms of transformations on distributed datasets*
- Resilient Distributed Datasets (RDDs)
 - Collections of objects that can be stored in memory or disk across a cluster
 - Parallel functional transformations (map, filter, ...)
 - Automatically rebuilt on failure

Spark for Data Science

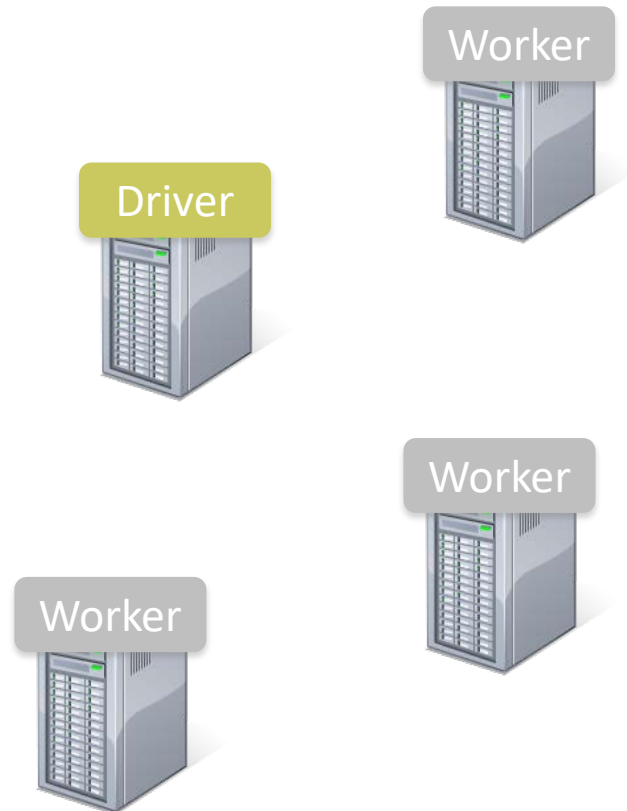
- DataFrames
 - Structured data
 - Familiar API based on R & Python Pandas
 - Distributed, optimized implementation
- Machine Learning Pipelines
 - Simple construction and tuning of ML workflows

Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns



Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
```



Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

Base RDD

```
lines = spark.textFile("hdfs://...")
```

Driver

Worker

Worker

Worker

Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))
```



Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

Transformed RDD

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))
```

Driver

Worker

Worker

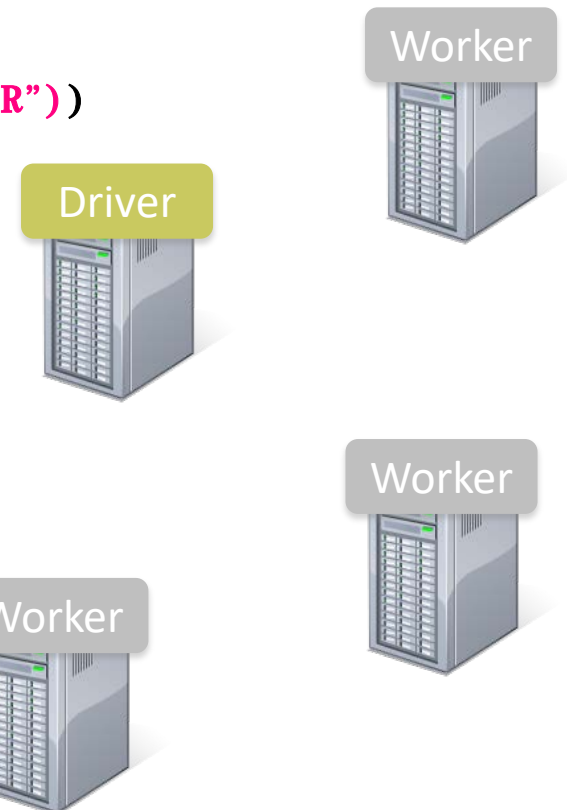
Worker

Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```



Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

Driver

Action

Worker

Worker

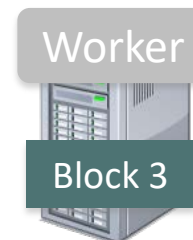
Worker

Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

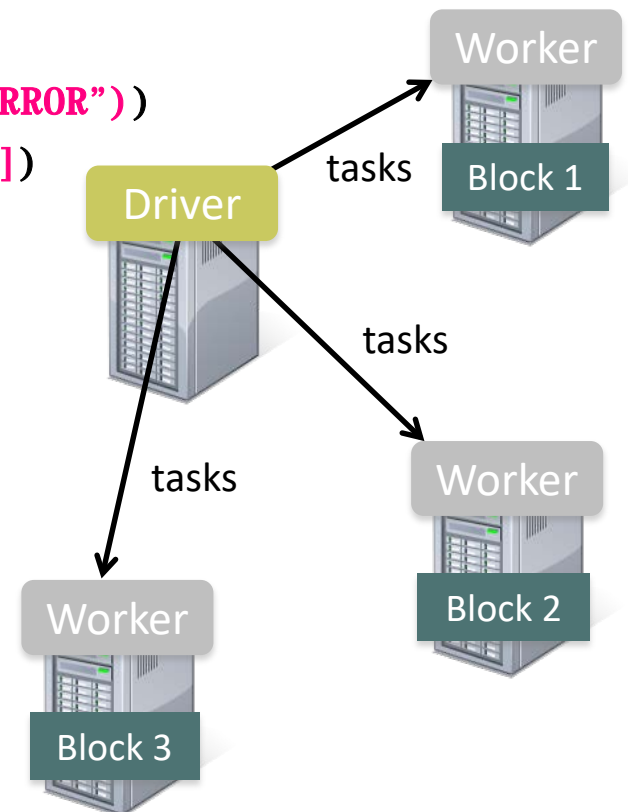


Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

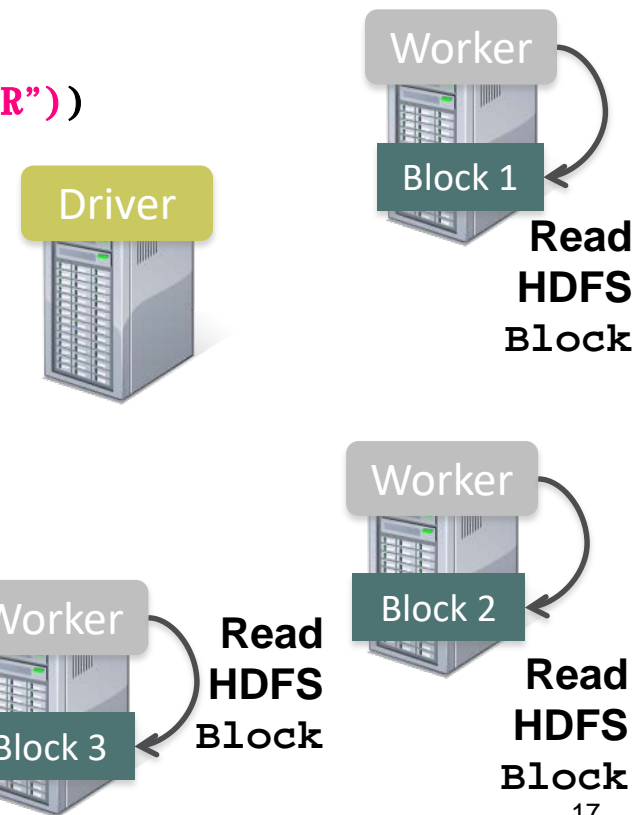


Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

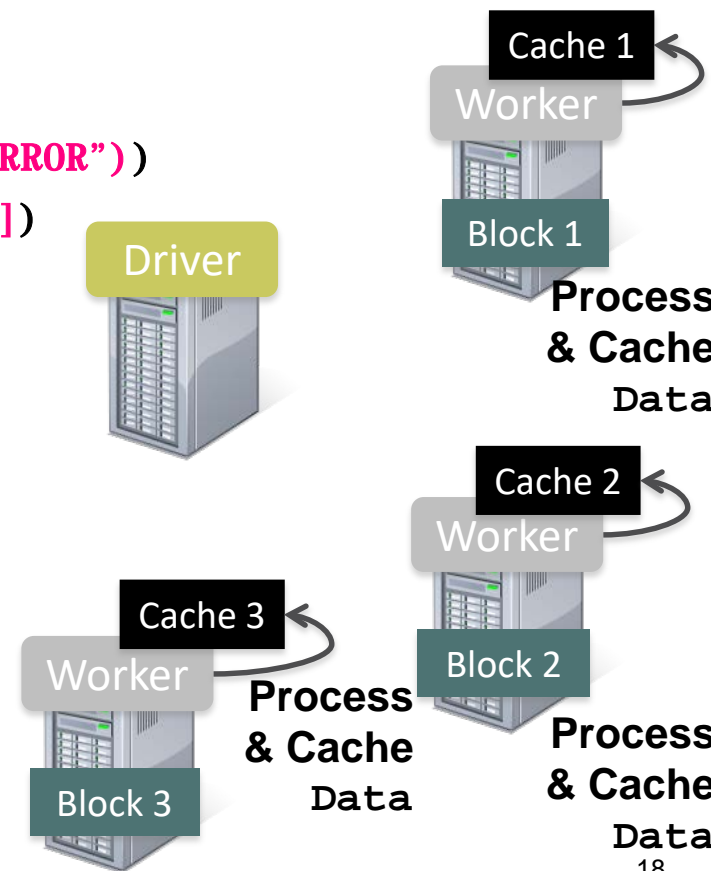


Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

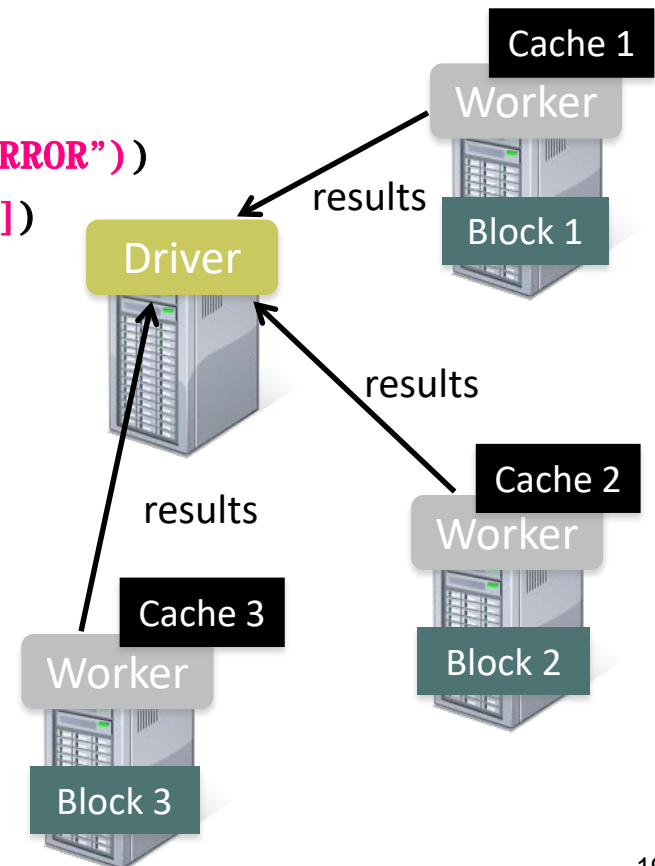


Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))  
messages = errors.map(lambda s: s.split("\t")[2])  
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

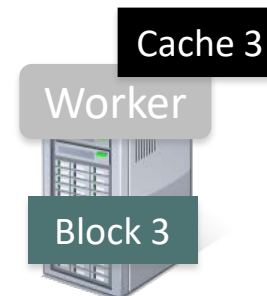
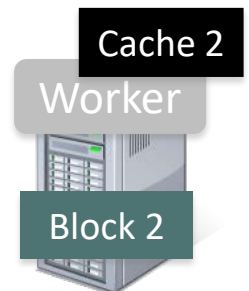
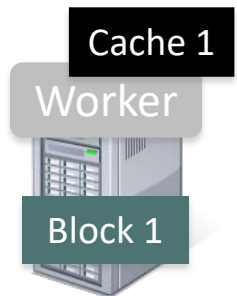


Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

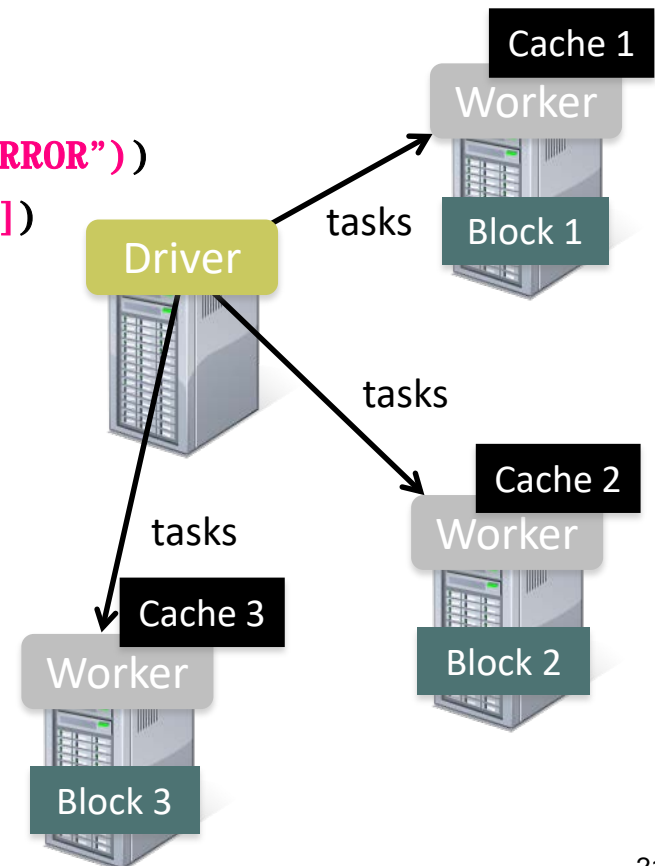


Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

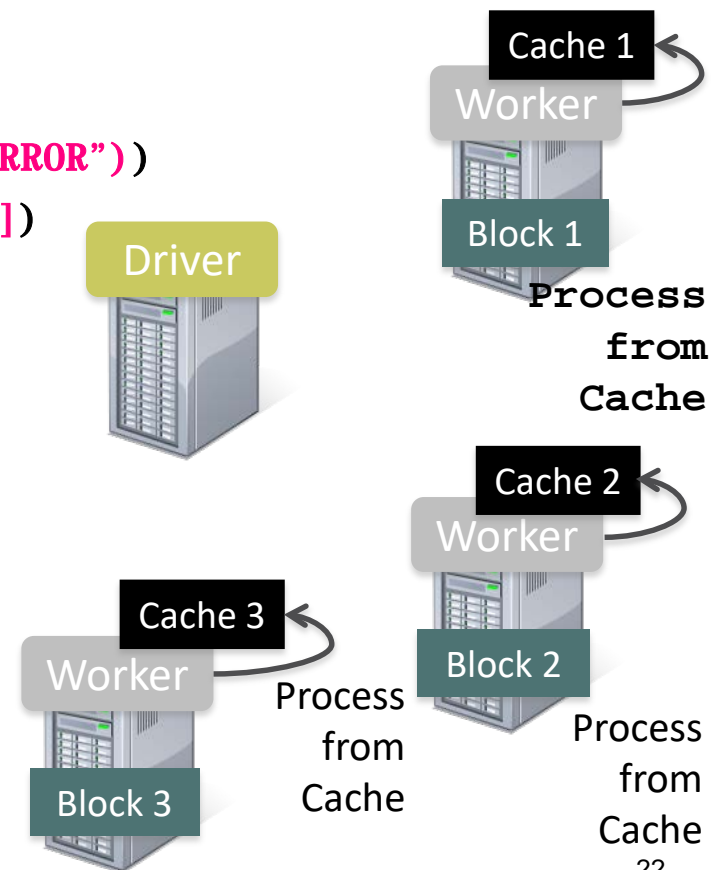


Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

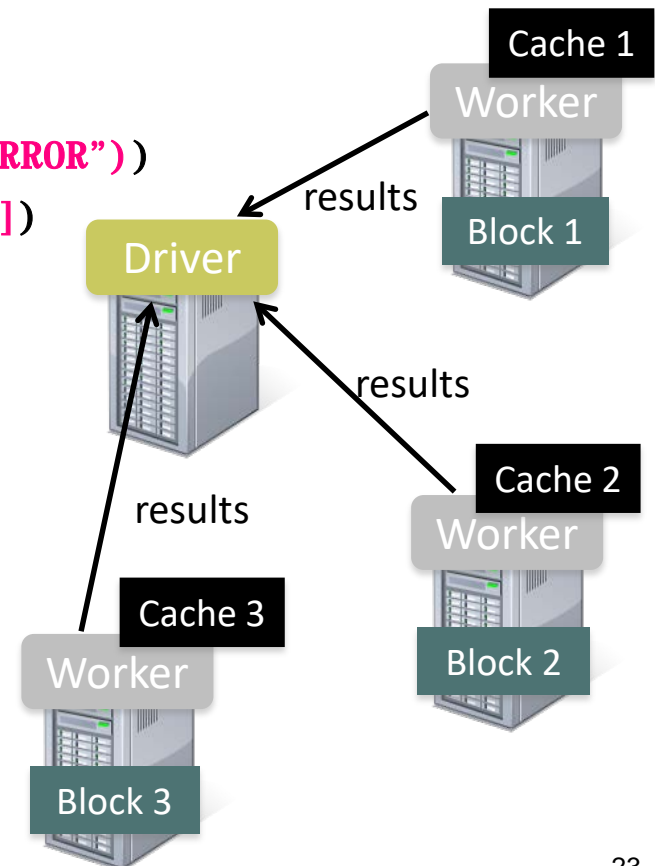


Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```



Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

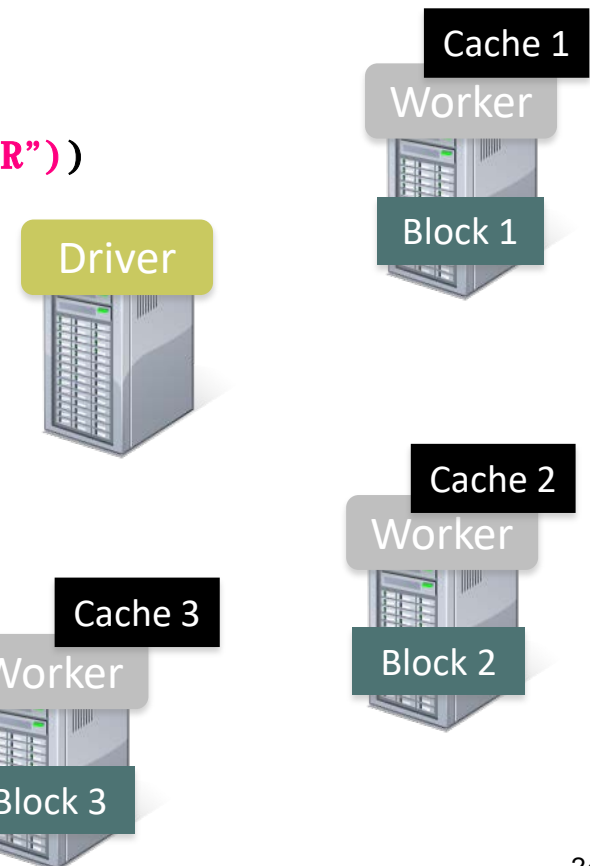
```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

Cache your data → Faster Results

Full-text search of Wikipedia

- 60GB on 20 EC2 machines
- 0.5 sec from mem vs. 20s for on-disk



Week 2 Contents

- Spark Recap – Example on Cache
- **Spark DataFrame & Dataset**
- Spark MLlib & ML Pipelines
- GitHub Classroom & Quiz 1

Challenges & Solutions

- Perform ETL to and from various (semi- or unstructured) data sources
- Perform advanced analytics (e.g. machine learning, graph processing) that are hard to express in relational systems
- A *DataFrame* API that can perform relational operations on both external data sources and Spark's built-in RDDs.
- A highly extensible optimizer, *Catalyst*, that uses features of Scala to add composable rule, control code gen., and define extensions.

DataFrame-based API for MLlib

- a.k.a. “Pipelines” API, with utilities for constructing ML Pipelines
- In 2.0, the DataFrame-based API will become the primary API for MLlib
 - Voted by community
 - `org.apache.spark.ml`, `pyspark.ml`
- The RDD-based API will enter maintenance mode
 - Still maintained with bug fixes, but no new features
 - `org.apache.spark.mllib`, `pyspark.mllib`

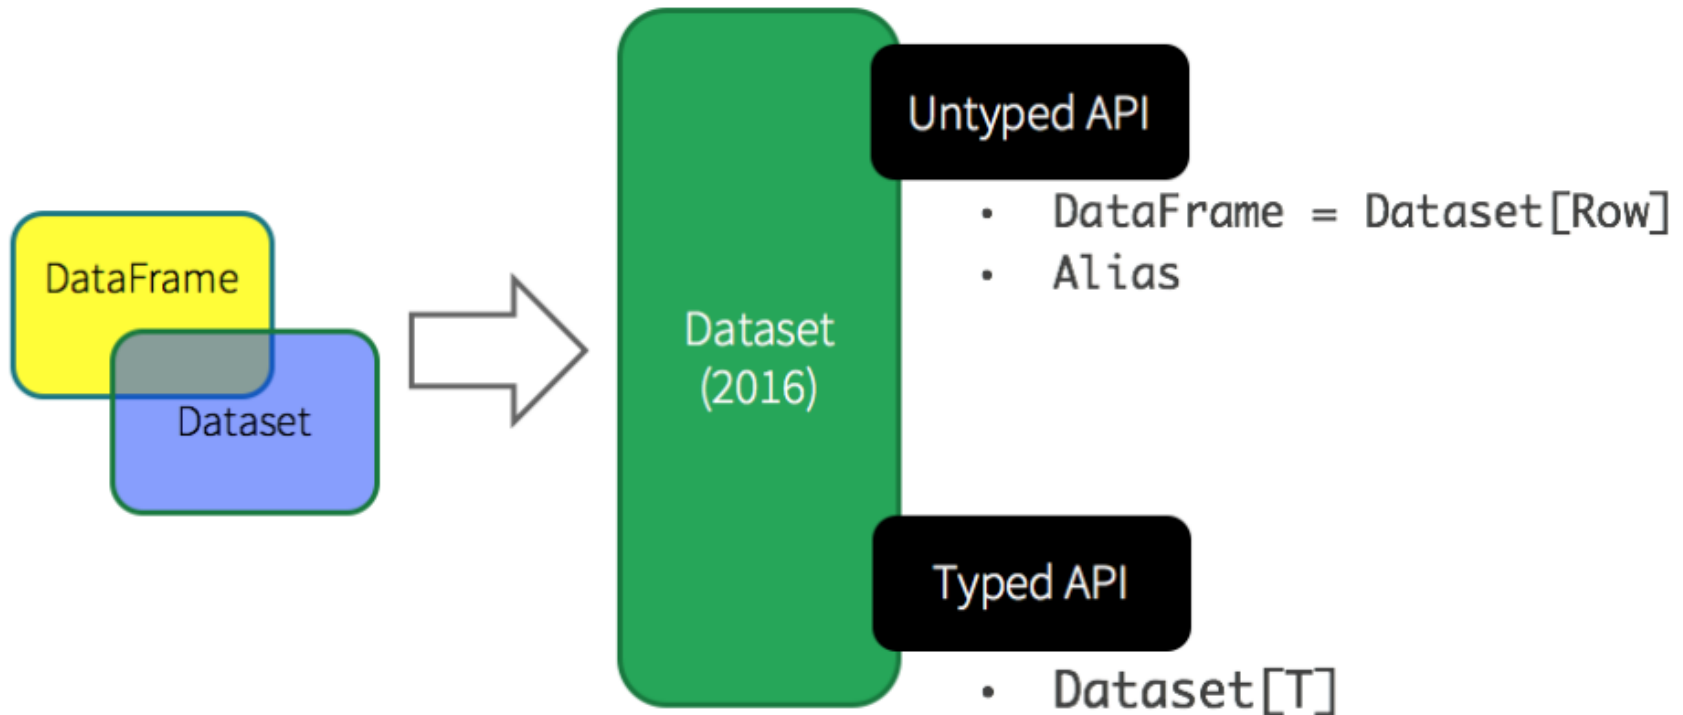
Structuring Spark: DataFrames and Datasets

- DataFrames and Datasets
 - DataFrame (schema, generic untyped)
 - Index access, named columns (like a table)
 - Dataset (static typing, strongly-typed)
 - Object access
 - DataFrame = Dataset[Row] (row: generic untyped)
 - Unified in Apache Spark 2.0
- RDDs are low-level (like assembler), DataFrames & Datasets are built on top of RDDs
- New libraries: built on Datasets and DataFrames

[*A Tale of Three Apache Spark APIs: RDDs, DataFrames, and Datasets*](#)

Unified API

Unified Apache Spark 2.0 API



Typed and Un-typed APIs

Language	Main Abstraction
Scala	Dataset[T] & DataFrame (alias for Dataset[Row])
Java	Dataset[T]
Python*	DataFrame
R*	DataFrame

Note: Since Python and R have no compile-time type-safety, we only have untyped APIs, namely DataFrames.

Benefits of Dataset APIs

- Static-typing and runtime type-safety
 - SQL least restrictive, no syntax error until runtime
 - DF/DS: syntax error detected at compile time
- High-level abstraction and custom view into structured and semi-structured data, e.g. JSON
- Ease-of-use of APIs with structure
 - Rich semantics and domain specific operations
- Performance and Optimization
 - SQL Catalyst

DataFrame

```
ctx = new HiveContext()
users = ctx.table("users")
young = users.where(users("age") < 21)
println(young.count())
```

- A distributed collection of rows with the same schema
- Can be constructed from external data sources or RDDs into essentially an RDD of Row objects
- Supports relational operators (e.g. *where*, *groupby*) as well as Spark operations.
- Evaluated lazily → unmaterialized *logical* plan

DataFrames

dept	age	name
Bio	48	H Smith
CS	54	A Turing
Bio	43	B Jones
Chem	61	M Kennedy

Data grouped into
named columns

RDD API

```
pdata.map(lambda x: (x.dept, [x.age, 1])) \  
      .reduceByKey(lambda x, y: [x[0] + y[0], x[1] + y[1]]) \  
      .map(lambda x: [x[0], x[1][0] / x[1][1]]) \  
      .collect()
```

DataFrame API

```
data.groupBy("dept").avg("age")
```

DataFrames

dept	age	name
Bio	48	H Smith
CS	54	A Turing
Bio	43	B Jones
Chem	61	M Kennedy

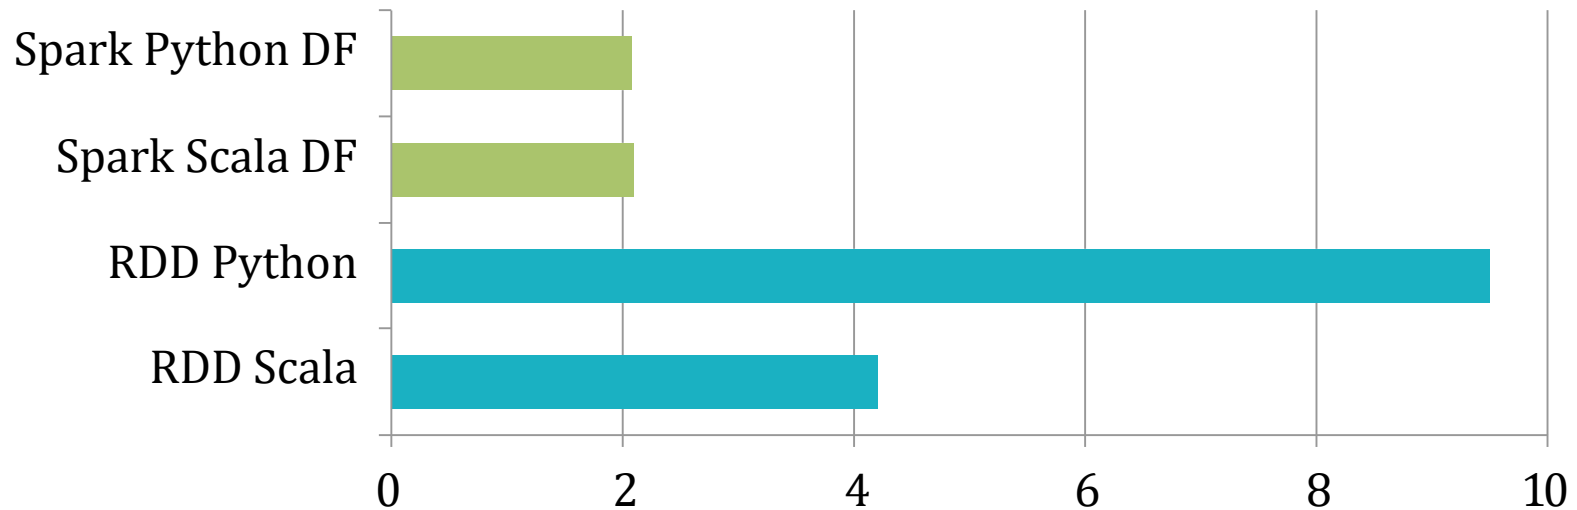
Data grouped into
named columns

DSL for common tasks

- Project, filter, aggregate, join, ...
- Metadata
- UDFs

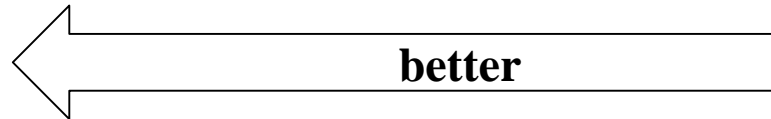
```
data.groupBy("dept").avg("age")
```

Spark DataFrames are *fast*



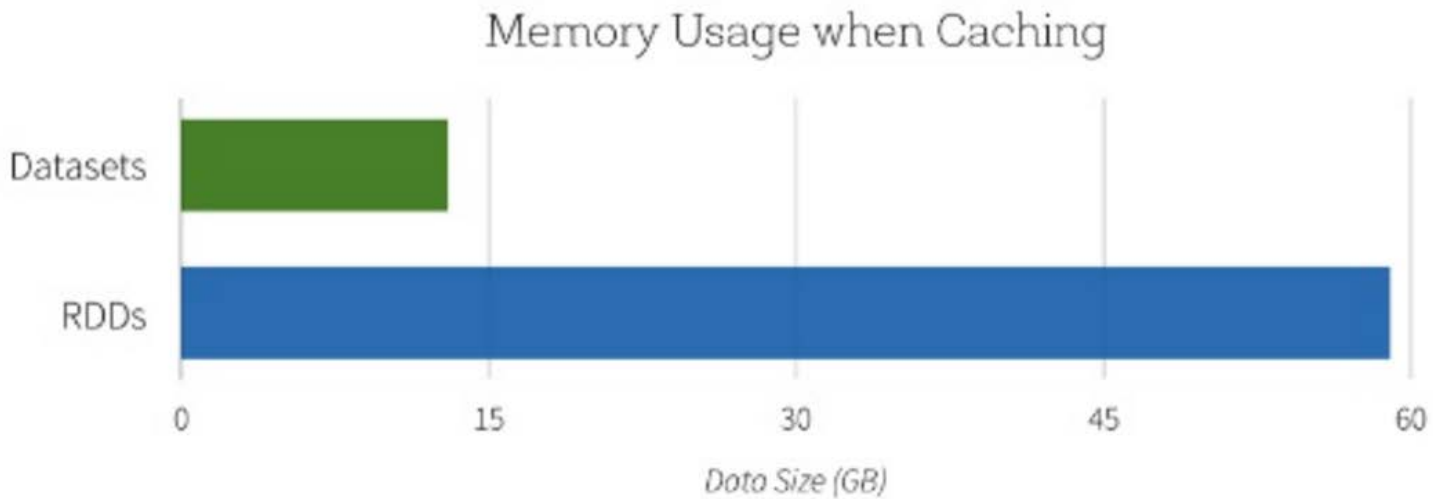
*Uses SparkSQL
Catalyst optimizer*

Runtime of aggregating 10 million int pairs (secs)



Space Efficiency

Space Efficiency



Week 2 Contents

- Spark Recap – Example on Cache
- Spark DataFrame & Dataset
- **Spark MLlib & ML Pipelines**
- GitHub Classroom & Quiz 1

Machine Learning Library (MLlib)

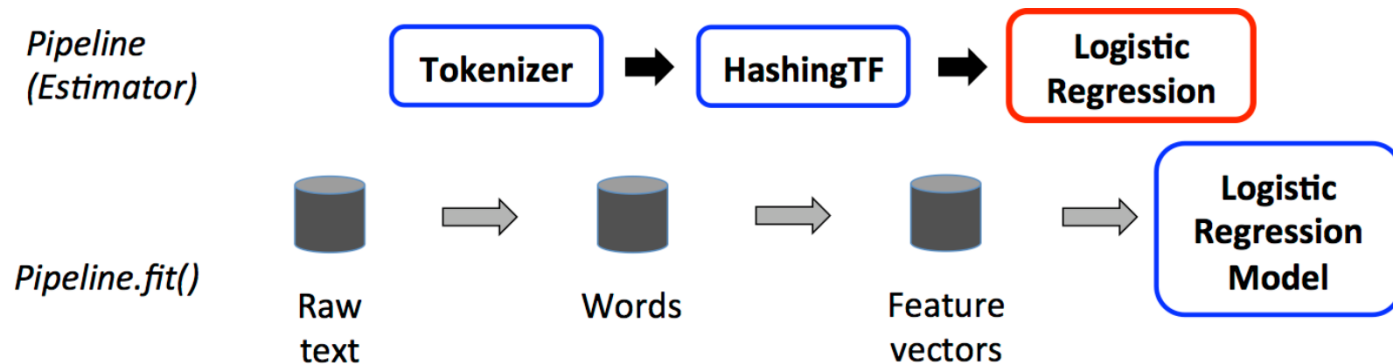
- ML Algorithms: common learning algorithms such as classification, regression, clustering, and collaborative filtering
- Featurization: feature extraction, transformation, dimensionality reduction, and selection
- Pipelines: tools for constructing, evaluating, and tuning ML Pipelines
- Persistence: saving and load algorithms, models, and Pipelines
- Utilities: linear algebra, statistics, data handling, etc

Main Concepts in Pipelines

- DataFrame: an ML dataset, which can hold a variety of data types. E.g., different columns storing text, feature vectors, true labels, and predictions.
- Transformer: algorithm transforming one DataFrame into another DataFrame. E.g., ML model → features into predictions.
- Estimator: algorithm fit on a DataFrame to produce a Transformer. E.g., ML algorithm DataFrame → model
- Pipeline: chains multiple Transformers and Estimators together to specify an ML workflow.
- Parameter: All Transformers and Estimators now share a common API for specifying parameters

ML Pipelines

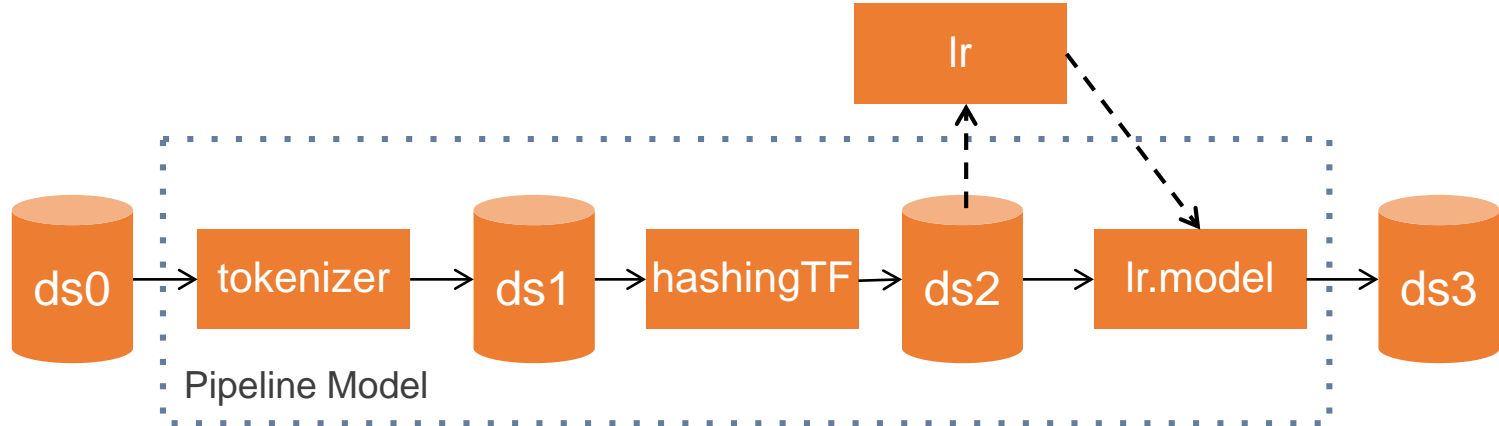
- ML Pipelines: high-level APIs to create and tune machine learning pipelines.
- Spark DataFrame: distributed collection of data organized into named columns.
 - Table in a relational database
 - Data frame in R or Python



Spark MLlib Pipelines

```
tokenizer = Tokenizer(inputCol="text", outputCol="words")
hashingTF = HashingTF(inputCol="words", outputCol="features")
lr = LogisticRegression(maxIter=10, regParam=0.01)
pipeline = Pipeline(stages=[tokenizer, hashingTF, lr])

df = sqlCtx.load("/path/to/data")
model = pipeline.fit(df)
```



Example: Text Classification

Goal: Given a text document, predict its topic.

Features

Subject: Re: Lexan Polish?
Suggest McQuires #1 plastic
polish. It will help somewhat
but nothing will remove deep
scratches without making it
worse than it already is.
McQuires will do
something...

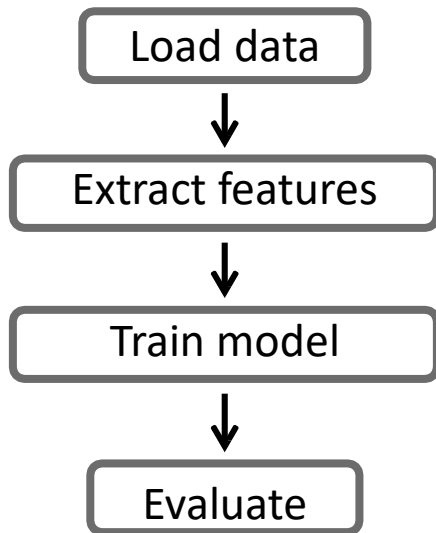


Label

1: about science
0: not about science

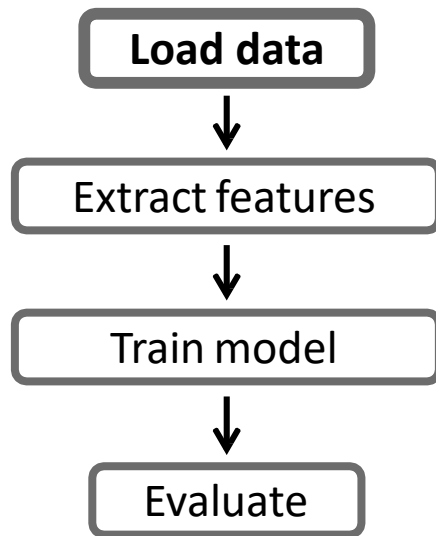
Dataset: "20 Newsgroups"
From UCI KDD Archive

ML Workflow



Load Data

Data sources for DataFrames



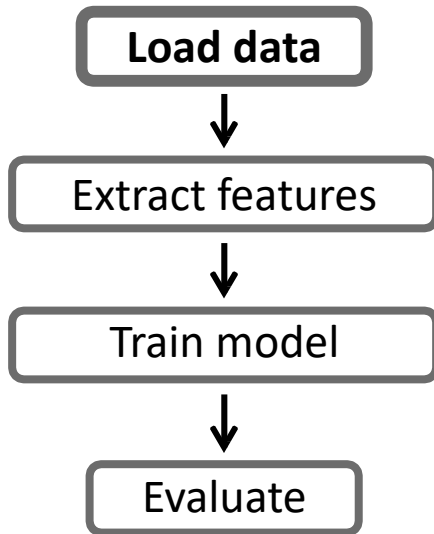
built-in



external



Load Data



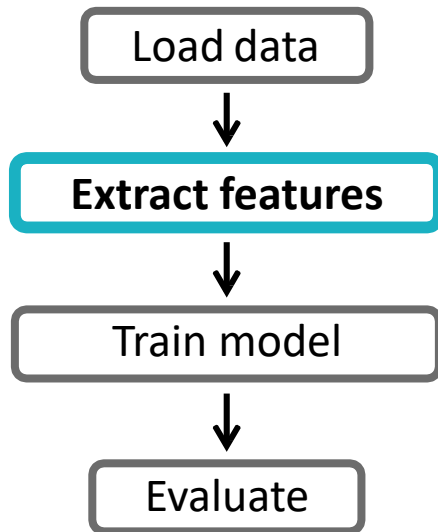
Current data schema

```
label: Int  
text: String
```

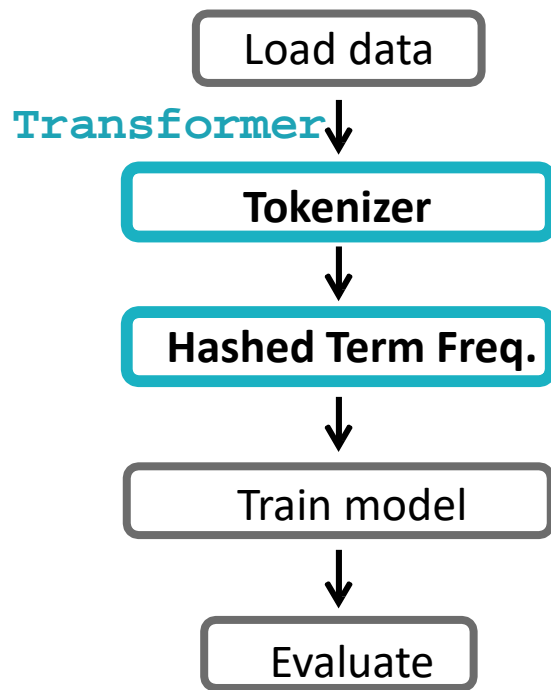
Extract Features

Current data schema

```
label: Int  
text: String
```



Extract Features



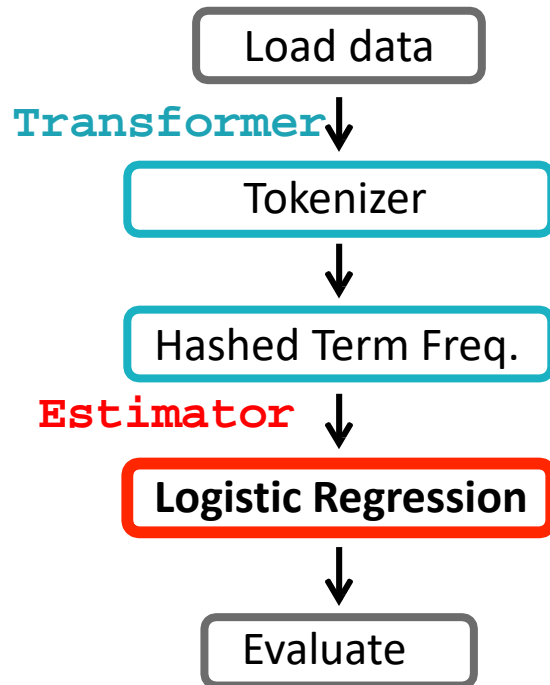
Current data schema

```
label: Int  
text: String
```

```
words: Seq[String]
```

```
features: Vector
```

Train a Model



Current data schema

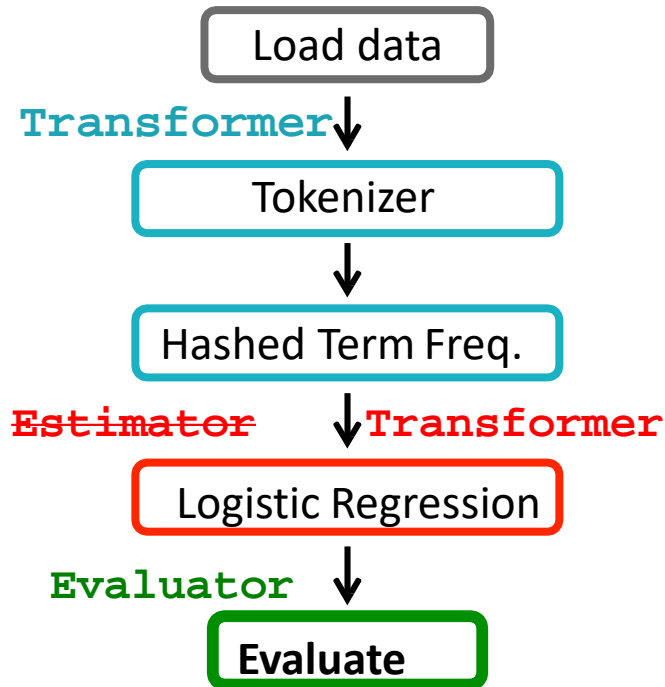
label: Int
text: String

words: Seq[String]

features: Vector

model parameters

Evaluate the Model



Current data schema

label: Int
text: String

words: Seq[String]

features: Vector

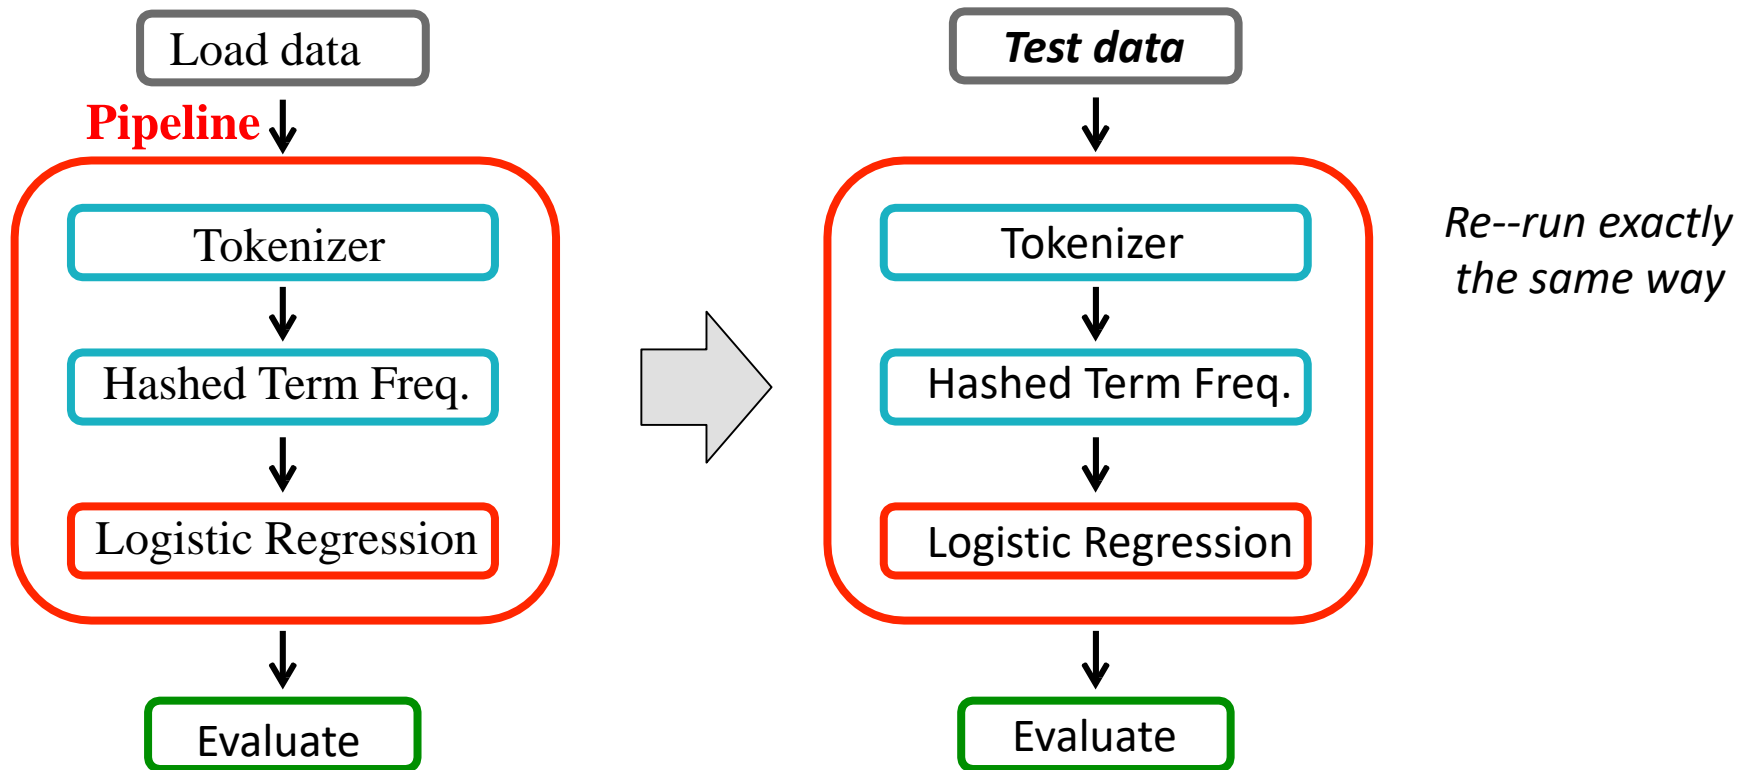
prediction: Int

By default, always append new columns

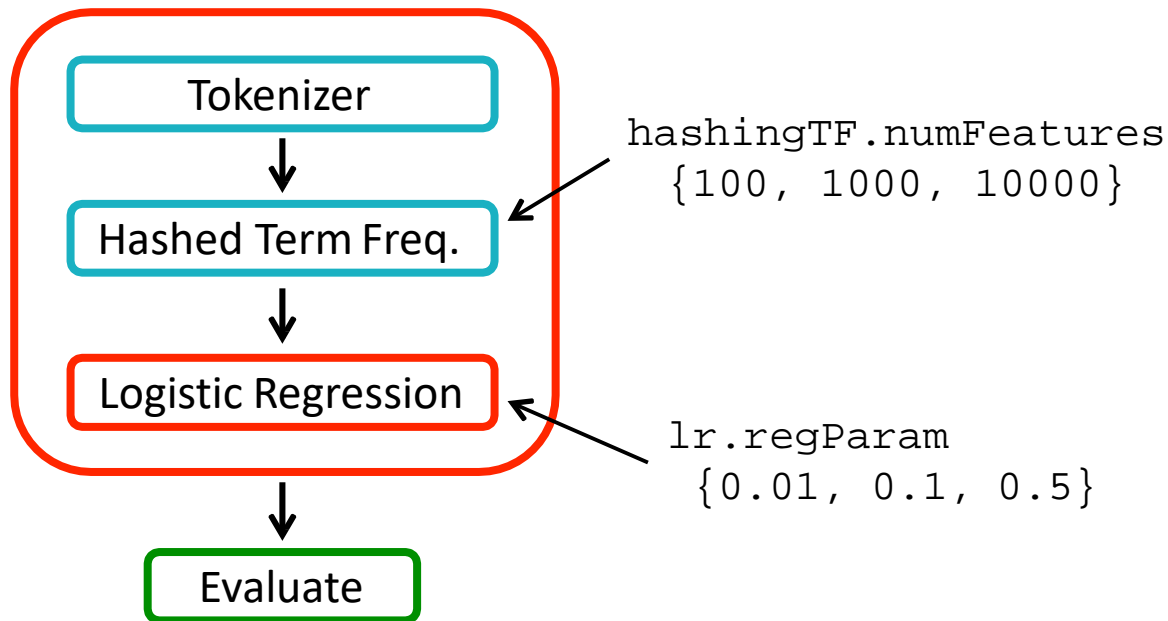
→ Can go back & inspect intermediate results

→ Made efficient by DataFrame optimizations

ML Pipelines



Parameter Tuning



CrossValidator

Given:

- Estimator
- Parameter grid
- Evaluator

Find best parameters

Week 2 Contents

- Spark Recap – Example on Cache
- Spark DataFrame & Dataset
- Spark MLlib & ML Pipelines
- **GitHub Classroom & Quiz 1**

GitHub Classroom & Quiz 1

- GitHub Classroom
 - Lab sessions
 - Quizzes
 - Assignments
 - [Feedbacks]
- **Quiz 1**
 - **22 Feb 2018 10am in lab session**
 - 50 minutes max
 - **10%** of your total mark

Recommended Reading

- Sections 2.4.2 and 2.4.3 of the MMDS book (3rd edition)

<http://i.stanford.edu/~ullman/mmds/ch2n.pdf>

- The full book
<http://i.stanford.edu/~ullman/mmds/book0n.pdf>