

# Lecture 8

## Scalable PCA/SVD

### *Dimensionality Reduction & Factor Analysis*

Haiping Lu

<http://www.dcs.shef.ac.uk/~haiping>

COM6012 Scalable Machine Learning

Spring 2018

# Week 8 Contents

- **Unsupervised Learning**
- PCA - Dimensionality Reduction
- SVD – Factor Analysis
- Scalable PCA in Spark

# Unsupervised Learning

Supervised methods

$$y = f(X)$$




predict  
our data

as a function of  
other data

**Unsupervised methods**

$$f(X)$$



**find structure in the  
data on its own**


# Three Topics

- Principal component analysis (PCA) & SVD
  - Dimensionality reduction & factor analysis
- K-means
  - Clustering
- Matrix factorisation (with missing information)
  - Collaborative filtering → Recommender system
- **Scale these algorithms for big data**

# Week 8 Contents

- Unsupervised Learning
- **PCA - Dimensionality Reduction**
- SVD – Factor Analysis
- Scalable PCA in Spark

# Dimensionality Reduction



- **Assumption:** Data lies on or near a low  $d$ -dimensional subspace
- Axes of this subspace are effective representation of the data

# Why Reduce Dimensions?

## **Why reduce dimensions?**

- Discover hidden correlations/topics
  - Words that occur commonly together
- Remove redundant and noisy features
  - Not all words are useful
- Interpretation and visualization
- Easier storage and processing of the data

# Dimensionality Reduction

- Raw data is complex and high-dimensional
- Dimensionality reduction describes the data using a simpler, more compact representation
- This representation may make interesting patterns in the data clearer or easier to see




# Dimensionality Reduction

- Goal: Find a ‘better’ representation for data
- How do we define ‘better’?
- For example
  - Minimise reconstruction error
  - Maximise variance
  - **They give the same solution → PCA!**

# PCA Algorithm


- Input:  $N$  data points, each  $\rightarrow D$ -dimensional vector
- PCA algorithm
  - 1.  $\mathbf{X}_0 \leftarrow$  Form  $N \times D$  data matrix, with one row vector  $\mathbf{x}_n$  per data point
  - 2.  $\mathbf{X}$ : subtract mean  $\mathbf{x}$  from each row vector  $\mathbf{x}_n$  in  $\mathbf{X}_0$
  - 3.  $\Sigma \leftarrow \mathbf{X}^T \mathbf{X}$  Gramian (scatter) matrix for  $\mathbf{X}$
  - Find eigenvectors and eigenvalues of  $\Sigma$
  - PCs  $\mathbf{U} (D \times d) \leftarrow$  the  $d$  eigenvectors with largest eigenvalues
- PCA feature for  $\mathbf{y}$   $D$ -dim:  $\mathbf{U}^T \mathbf{y}$  ( $d$ -dimensional)
  - Zero correlations, ordered by variance

# 2D Data




# Principal Components

- The best axis to project
- Minimum RMS error
- Principal vectors are **orthogonal**



# How Many Components?

- Check the distribution of eigen-values
- Take enough many eigen-vectors to cover 80-90% of the variance



# Other Practical Tips

- PCA assumptions (linearity, orthogonality) not always appropriate
- Various extensions to PCA with different underlying assumptions, e.g., manifold learning, Kernel PCA, ICA
- Centring is crucial, i.e., we must preprocess data so that all features have zero mean before applying PCA
- PCA results dependent on scaling of data
- Data is sometimes rescaled in practice before applying PCA

# Problems and Limitations

- What if very large dimensional data?
  - e.g., Images ( $D \geq 10^4 = 100 \times 100$ )
- Problem:
  - Gramian matrix  $\Sigma$  is size ( $D^2$ )
  - $D=10^4 \rightarrow |\Sigma| = 10^8$
- Singular Value Decomposition (SVD)!
  - Efficient algorithms available
  - Some implementations find just top  $d$  eigenvectors

# Week 8 Contents

- Unsupervised Learning
- PCA - Dimensionality Reduction
- **SVD – Factor Analysis**
- Scalable PCA in Spark



# Singular Value Decomposition

- Factorization (decomposition) problem
  - #1: Find concepts/topics/genres → Factor Analysis
  - #2: Reduce dimensionality

term	data	information	retrieval	brain	lung
document					
CS-TR1	1	1	1	0	0
CS-TR2	2	2	2	0	0
CS-TR3	1	1	1	0	0
CS-TR4	5	5	5	0	0
MED-TR1	0	0	0	2	2
MED-TR2	0	0	0	3	3
MED-TR3	0	0	0	1	1

The above matrix is actually “2-dimensional.” All rows can be reconstructed by scaling  $[1 \ 1 \ 1 \ 0 \ 0]$  or  $[0 \ 0 \ 0 \ 1 \ 1]$ :  $D=5 \rightarrow d=2$

# SVD - Definition

$$\mathbf{A}_{[n \times m]} = \mathbf{U}_{[n \times r]} \mathbf{\Lambda}_{[r \times r]} (\mathbf{V}_{[m \times r]})^T$$

- $\mathbf{A}$ :  $n \times m$  matrix (e.g.,  $n$  documents,  $m$  terms)
- $\mathbf{U}$ :  $n \times r$  matrix ( $n$  documents,  $r$  concepts)
- $\mathbf{\Lambda}$ :  $r \times r$  diagonal matrix (strength of each ‘concept’) ( $r$ : rank of the matrix)
- $\mathbf{V}$ :  $m \times r$  matrix ( $m$  terms,  $r$  concepts)

# SVD - Properties

Always possible to decompose matrix  $\mathbf{A}$  into  $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$ ,  
where

- $\mathbf{U}, \mathbf{\Lambda}, \mathbf{V}$ : unique (\*)
- $\mathbf{U}, \mathbf{V}$ : column orthonormal (i.e., columns are unit vectors, orthogonal to each other)
  - $\mathbf{U}^T \mathbf{U} = \mathbf{I}; \mathbf{V}^T \mathbf{V} = \mathbf{I}$  ( $\mathbf{I}$ : identity matrix)
- $\mathbf{\Lambda}$ : singular value are positive, and sorted in decreasing order

# SVD $\leftrightarrow$ Eigen-decomposition

- SVD gives us:
  - $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$
- Eigen-decomposition:
  - $\mathbf{B} = \mathbf{W} \mathbf{\Sigma} \mathbf{W}^T$ 
    - $\mathbf{U}, \mathbf{V}, \mathbf{W}$  are orthonormal ( $\mathbf{U}^T \mathbf{U} = \mathbf{I}$ ),
    - $\mathbf{\Lambda}, \mathbf{\Sigma}$  are diagonal
- Relationship:
  - $\mathbf{A} \mathbf{A}^T = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T (\mathbf{U} \mathbf{\Lambda} \mathbf{V}^T)^T = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T (\mathbf{V} \mathbf{\Lambda}^T \mathbf{U}^T) = \mathbf{U} \mathbf{\Lambda} \mathbf{\Lambda}^T \mathbf{U}^T$
  - $\mathbf{A}^T \mathbf{A} = \mathbf{V} \mathbf{\Lambda}^T \mathbf{U}^T (\mathbf{U} \mathbf{\Lambda} \mathbf{V}^T) = \mathbf{V} \mathbf{\Lambda} \mathbf{\Lambda}^T \mathbf{V}^T = \mathbf{V} \mathbf{\Lambda}^2 \mathbf{V}^T$
  - $\mathbf{B} = \mathbf{A}^T \mathbf{A} = \mathbf{W} \mathbf{\Sigma} \mathbf{W}^T$

# SVD for PCA

- PCA by SVD:
  - 1.  $\mathbf{X}_0 \leftarrow$  Form  $N \times d$  data matrix, with one row vector  $\mathbf{x}_n$  per data point
  - 2.  $\mathbf{X}$  subtract mean  $\mathbf{x}$  from each row vector  $\mathbf{x}_n$  in  $\mathbf{X}_0$
  - 3.  $\mathbf{U} \mathbf{\Lambda} \mathbf{V}^T \leftarrow$  SVD of  $\mathbf{X}$
  - The right singular vectors  $\mathbf{V}$  of  $\mathbf{X}$  are equivalent to the eigenvectors of  $\mathbf{X}^T \mathbf{X} \rightarrow$  the PCs
  - The singular values in  $\mathbf{\Lambda}$  are equal to the square roots of the eigenvalues of  $\mathbf{X}^T \mathbf{X}$

# SVD - Properties

‘spectral decomposition’ of the matrix:

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} | & | \\ \mathbf{u}_1 & \mathbf{u}_2 \\ | & | \end{bmatrix} \mathbf{X} \begin{bmatrix} \lambda_1 & \emptyset \\ \emptyset & \lambda_2 \end{bmatrix} \mathbf{X} \begin{bmatrix} \text{---} \mathbf{v}_1 \text{---} \\ \text{---} \mathbf{v}_2 \text{---} \end{bmatrix}$$

# SVD - Interpretation

‘documents’, ‘terms’ and ‘concepts’:

- $U$ : document-to-concept similarity matrix
- $V$ : term-to-concept similarity matrix
- $\Lambda$ : its diagonal elements: ‘strength’ of each concept

Projection:

- Best axis to project on: (‘best’ = min sum of squares of projection errors)

# SVD - Example

- $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$  - example:

term	data	information	retrieval	brain	lung
document					
CS-TR1	1	1	1	0	0
CS-TR2	2	2	2	0	0
CS-TR3	1	1	1	0	0
CS-TR4	5	5	5	0	0
MED-TR1	0	0	0	2	2
MED-TR2	0	0	0	3	3
MED-TR3	0	0	0	1	1

$$\begin{array}{c}
 \uparrow \\
 \text{CS} \\
 \downarrow \\
 \uparrow \\
 \text{MD} \\
 \downarrow
 \end{array}
 \begin{bmatrix}
 1 & 1 & 1 & 0 & 0 \\
 2 & 2 & 2 & 0 & 0 \\
 1 & 1 & 1 & 0 & 0 \\
 5 & 5 & 5 & 0 & 0 \\
 0 & 0 & 0 & 2 & 2 \\
 0 & 0 & 0 & 3 & 3 \\
 0 & 0 & 0 & 1 & 1
 \end{bmatrix}
 =
 \begin{bmatrix}
 0.18 & 0 \\
 0.36 & 0 \\
 0.18 & 0 \\
 0.90 & 0 \\
 0 & 0.53 \\
 0 & 0.80 \\
 0 & 0.27
 \end{bmatrix}
 \times
 \begin{bmatrix}
 9.64 & 0 \\
 0 & 5.29
 \end{bmatrix}
 \times
 \begin{bmatrix}
 0.58 & 0.58 & 0.58 & 0 & 0 \\
 0 & 0 & 0 & 0.71 & 0.71
 \end{bmatrix}$$



# SVD - Example

- $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$  - example:

doc-to-concept  
similarity matrix

CS-concept  
MD-concept

↑

CS

↓

↑

MD

↓

		retrieval				
	data	inf. ↓	brain	lung		
	↑					
	↓					
	↑					
	↓					
	↑					
	↓					
	↑					
	↓					

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

*Note: In the original image, the value 0.18 in the first row, first column of the middle matrix is circled in blue, with arrows pointing to it from the labels 'CS-concept' and 'MD-concept'.*

# SVD - Example

- $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$  - example:

retrieval  
inf. ↓    brain    lung

data    'strength' of CS-concept

↑

CS

↓

↑

MD

↓

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix}$$
 $=$ 

$$\begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix}$$

$\times$ 

$$\begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix}$$
 $\times$

$$\begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

# SVD - Example

- $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$  - example:

retrieval

inf. ↓

data    brain    lung

↑ CS

↓

↑ MD

↓

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

term-to-concept  
similarity matrix

CS-concept

# SVD – Dimensionality Reduction

- Q: how exactly is (**further**) dim. reduction done?
- A: set the smallest singular values to zero:
- Note: **3 zero singular values** already removed

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

Diagram illustrating the SVD decomposition of a 7x5 matrix. The original matrix is decomposed into three matrices: a 7x2 matrix of left singular vectors, a 2x2 diagonal matrix of singular values, and a 2x5 matrix of right singular vectors. The singular values 9.64 and 5.29 are shown, with the value 5.29 crossed out by a blue 'X', indicating it is the smallest singular value and thus the target for dimensionality reduction. The right singular vector matrix shows the first two columns (0.58, 0.58, 0.58, 0, 0) and the last two columns (0, 0, 0, 0.71, 0.71), with the last two columns crossed out by a red 'X', indicating they are the smallest singular vectors and thus the target for dimensionality reduction.

# SVD - Dimensionality Reduction

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} \sim \begin{bmatrix} 0.18 \\ 0.36 \\ 0.18 \\ 0.90 \\ 0 \\ 0 \\ 0 \end{bmatrix} \times \begin{bmatrix} 9.64 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \end{bmatrix}$$

# SVD - Dimensionality Reduction

- Best rank-1 approximation

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} \sim \begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

# Week 8 Contents

- Unsupervised Learning
- PCA - Dimensionality Reduction
- SVD – Factor Analysis
- **Scalable PCA in Spark**

# PCA & SVD in Spark MLlib

- Not scalable: `computePrincipalComponents( )` from `RowMatrix`
- **Scalable:** `computeSVD( )` from `RowMatrix`
- Code:  
<https://github.com/apache/spark/blob/v2.1.0/mllib/src/main/scala/org/apache/spark/mllib/linalg/distributed/RowMatrix.scala>
- Documentation:  
<https://spark.apache.org/docs/2.1.0/api/scala/index.html#org.apache.spark.mllib.linalg.distributed.RowMatrix>



# PCA in Spark MLlib (RDD)

- <https://spark.apache.org/docs/2.1.0/mllib-dimensionality-reduction.html>

```
val mat: RowMatrix = new RowMatrix(dataRDD)

// Compute the top 4 principal components.
// Principal components are stored in a local dense matrix.
val pc: Matrix = mat.computePrincipalComponents(4)
```

- Not scalable, local computation

```
val brzSvd.SVD(u: BDM[Double], s: BDV[Double], _) = brzSvd(Cov)
```

- Notebook 8

# PCA in Spark ML (DF)

- Now in

<https://spark.apache.org/docs/2.1.0/ml-features.html#pca>

- Under features
- Scalable? Not likely

```
val pca = new PCA()  
  .setInputCol("features")  
  .setOutputCol("pcaFeatures")  
  .setK(3)  
  .fit(df)
```

# SVD in Spark MLlib (RDD)

- <https://spark.apache.org/docs/2.1.0/mllib-dimensionality-reduction.html>
- With distributed implementations

```
val mat: RowMatrix = new RowMatrix(dataRDD)

// Compute the top 5 singular values and corresponding singular vectors.
val svd: SingularValueDecomposition[RowMatrix, Matrix] = mat.computeSVD(5, computeU = true)
val U: RowMatrix = svd.U // The U factor is a RowMatrix.
val s: Vector = svd.s // The singular values are stored in a local dense vector.
val V: Matrix = svd.V // The V factor is a local dense matrix.
```

# SVD in Spark MLlib (RDD)

- An  $m \times n$  data matrix  $\mathbf{A}$  with  $m > n$  (note different notations)
- For large matrices, usually we don't need the complete factorization but only the top  $k$  singular values and its associated singular vectors.
- Save storage, de-noise and recover the low-rank structure of the matrix (dimensionality reduction)

# SVD in Spark MLlib (RDD)

- An  $m \times n$  data matrix  $\mathbf{A}$
- Assume  $m > n$ , SVD  $\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$
- The singular values and the right singular vectors are derived from the eigenvalues and the eigenvectors of  $\mathbf{A}^T \mathbf{A}$  (which is smaller than  $\mathbf{A}$ )
- The left singular vectors are computed via matrix multiplication as  $\mathbf{U} = \mathbf{A} \mathbf{V} \mathbf{\Lambda}^{-1}$ , if requested by the user via the computeU parameter

# Selection of SVD Computation

- Auto
- If  $n$  is small ( $n < 100$ ) or  $k$  is large compared with  $n$  ( $k > n/2$ ), compute  $\mathbf{A}^T \mathbf{A}$  first and then compute its top eigenvalues and eigenvectors **locally** on the driver
- Otherwise, compute  $\mathbf{A}^T \mathbf{A} \mathbf{v}$  in a distributive way and send it to ARPACK to compute the top eigenvalues and eigenvectors on the driver node

# Selection of SVD Computation

- Auto (default)

```
if (n < 100 || (k > n / 2 && n <= 15000)) {  
    // If n is small or k is large compared with n, we better compute the Gramian matrix first  
    // and then compute its eigenvalues locally, instead of making multiple passes.  
    if (k < n / 3) {  
        SVDMode.LocalARPACK  
    } else {  
        SVDMode.LocalLAPACK  
    }  
} else {  
    // If k is small compared with n, we use ARPACK with distributed multiplication.  
    SVDMode.DistARPACK  
}
```

# Selection of SVD Computation

- Specify computeMode (private)

```
case "local-svd" => SVDMODE.LocalLAPACK  
case "local-eigs" => SVDMODE.LocalARPACK  
case "dist-eigs" => SVDMODE.DistARPACK
```



# Selection of SVD Computation

- computeMode (note brzSvd.SVD is local)

```
// Compute the eigen-decomposition of A' * A.
val (sigmaSquares: BDV[Double], u: BDM[Double]) = computeMode match {
  case SVDMode.LocalARPACK =>
    require(k < n, s"k must be smaller than n in local-eigs mode but got k=$k and n=$n.")
    val G = computeGramianMatrix().asBreeze.asInstanceOf[BDM[Double]]
    EigenValueDecomposition.symmetricEigs(v => G * v, n, k, tol, maxIter)
  case SVDMode.LocalLAPACK =>
    // breeze (v0.10) svd latent constraint, 7 * n * n + 4 * n < Int.MaxValue
    require(n < 17515, s"$n exceeds the breeze svd capability")
    val G = computeGramianMatrix().asBreeze.asInstanceOf[BDM[Double]]
    val brzSvd.SVD(uFull: BDM[Double], sigmaSquaresFull: BDV[Double], _) = brzSvd(G)
    (sigmaSquaresFull, uFull)
  case SVDMode.DistARPACK =>
    if (rows.getStorageLevel == StorageLevel.NONE) {
      logWarning("The input data is not directly cached, which may hurt performance if its"
        + " parent RDDs are also uncached.")
    }
    require(k < n, s"k must be smaller than n in dist-eigs mode but got k=$k and n=$n.")
    EigenValueDecomposition.symmetricEigs(multiplyGramianMatrixBy, n, k, tol, maxIter)
}
```

# Remark

- Acknowledgement
  - Some slides are adapted from slides by Jure Leskovec et al. <http://www.mmds.org>
- References
  - <http://infolab.stanford.edu/~ullman/mmds/ch11.pdf>
  - <http://www.mmds.org>
  - [https://en.wikipedia.org/wiki/Principal\\_component\\_analysis](https://en.wikipedia.org/wiki/Principal_component_analysis)
  - <https://spark.apache.org/docs/2.1.0/mllib-dimensionality-reduction.html>